

计算机科学与技术学院神经网络与深度学习课程实验报告

实验题目：卷积神经网络构建及应用		学号：201800130086
日期：2020.11.16	班级：18 智能班	姓名：徐鹏博
Email：hsupengbo@163.com		

实验目的：

理解卷积神经网络，构建神经网络模型并应用

- 完成 Convolutional model -Step by Step
- 完成 Convolutional model - Application

实验软件和硬件环境：

Anaconda3-64bit

Tensorflow1.15

Python 3.6

实验原理和方法：

我们要实现一个拥有卷积层（CONV）和池化层（POOL）的网络，它包含了前向和反向传播。

卷积模块，包含函数：使用 0 扩充边界，卷积窗口，前向卷积，反向卷积（可选）

池化模块，包含函数：前向池化，创建掩码，值分配，反向池化（可选）

卷积神经网络 - 前向传播

在前向传播的过程中，我们将使用多种过滤器对输入的数据进行卷积操作，每个过滤器会产生一个 2D 的矩阵，我们可以把它们堆叠起来，于是这些 2D 的卷积矩阵就变成了高维的矩阵。

$$n_H = \lfloor \frac{n_{H_{prev}} - f + 2 \times pad}{stride} \rfloor + 1$$

$$n_W = \lfloor \frac{n_{W_{prev}} - f + 2 \times pad}{stride} \rfloor + 1$$

$$n_C = \text{过滤器数量}$$

池化层的前向传播

要在同一个函数中实现最大值池化层和均值池化层

$$n_H = \lfloor \frac{n_{H_{prev}} - f}{stride} \rfloor + 1$$

$$n_W = \lfloor \frac{n_{W_{prev}} - f}{stride} \rfloor + 1$$

$$n_C = n_{C_{prev}}$$

最大值池化层：在输入矩阵中滑动一个大小为 $f \times f$ 的窗口，选取窗口里的值中的最大值，然后作为输出的一部分。

均值池化层：在输入矩阵中滑动一个大小为 $f \times f$ 的窗口，计算窗口里的值中的平均值，然后这个均值作为输出的一部分。

卷积层的反向传播

$$dA+ = \sum_{h=0}^{n_H} \sum_{w=0}^{n_W} W_c \times dZ_{hw}$$

$$dW_c+ = \sum_{h=0}^{n_H} \sum_{w=0}^{n_W} a_{slice} \times dZ_{hw}$$

$$db = \sum_h \sum_w dZ_{hw}$$

实验步骤：（不要求罗列完整源代码）

Convolutional Neural Networks: Step by Step

Convolutional Neural Networks

Zero-Padding

```
def zero_pad(X, pad):
    X_pad = np.pad(X, ((0,0),(pad,pad),(pad,pad),(0,0)), 'constant')
    return X_pad
```

Single step of convolution

```
def conv_single_step(a_slice_prev, W, b):
    s = a_slice_prev * W + b
    Z = np.sum(s)
    return Z
```

Convolutional Neural Networks - Forward pass

```
def conv_forward(A_prev, W, b, hparameters):
    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
    (f, f, n_C_prev, n_C) = W.shape
    stride = hparameters['stride']
    pad = hparameters['pad']
    n_H = int((n_H_prev - f + 2*pad)/stride + 1)
    n_W = int((n_W_prev - f + 2*pad)/stride + 1)
    Z = np.zeros((m,n_H,n_W,n_C))
    A_prev_pad = zero_pad(A_prev, pad)
    for i in range(m): # loop over the batch of training examples
        a_prev_pad = A_prev_pad[i,:,:,:]
        for h in range(n_H): # loop over vertical axis of the output volume
            for w in range(n_W): # loop over horizontal axis of the output volume
                for c in range(n_C): # loop over channels (= #filters) of the output volume
                    # Compute the slice from a_prev_pad (start=, end=)
                    a_slice_prev = a_prev_pad[h*hStride:(h+1)*hStride, w*wStride:(w+1)*wStride, :]
                    # Compute the output neuron (of shape 1x1x1)
                    Z[i,h,w,c] = conv_single_step(a_slice_prev, W[:,:,:,c], b[c])
```

```

vert_start = h * stride
vert_end = vert_start + stride
horiz_start = w * stride
horiz_end = horiz_start + stride
a_slice_prev = a_prev_pad[vert_start:vert_end,horiz_start:horiz_end,:]
Z[i, h, w, c] = conv_single_step(a_slice_prev, W[:, :, :, c], b[:, :, :, c])
assert(Z.shape == (m, n_H, n_W, n_C))
cache = (A_prev, W, b, hparameters)
return Z, cache

```

Pooling layer

Forward Pooling

```

def pool_forward(A_prev, hparameters, mode = "max"):
    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
    f = hparameters["f"]
    stride = hparameters["stride"]
    n_H = int(1 + (n_H_prev - f) / stride)
    n_W = int(1 + (n_W_prev - f) / stride)
    n_C = n_C_prev
    A = np.zeros((m, n_H, n_W, n_C))

    for i in range(m):           # loop over the training examples
        for h in range(n_H):      # loop on the vertical axis of the output volume
            for w in range(n_W):  # loop on the horizontal axis of the output volume
                for c in range (n_C): # loop over the channels of the output volume

                    vert_start = h * stride
                    vert_end = vert_start+f
                    horiz_start = w * stride
                    horiz_end = horiz_start + f
                    a_prev_slice = A_prev[i,vert_start:vert_end,horiz_start:horiz_end,c]
                    if mode == "max":
                        A[i, h, w, c] = np.max(a_prev_slice)
                    elif mode == "average":
                        A[i, h, w, c] = np.mean(a_prev_slice)

    cache = (A_prev, hparameters)
    assert(A.shape == (m, n_H, n_W, n_C))
    return A, cache

```

Backpropagation in convolutional neural networks

Convolutional layer backward pass

```

def conv_backward(dZ, cache):
    (A_prev, W, b, hparameters) = cache
    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
    (f, f, n_C_prev, n_C) = W.shape

```

```

stride = hparameters['stride']
pad = hparameters['pad']
(m, n_H, n_W, n_C) = dZ.shape

dA_prev = np.zeros((m,n_H_prev,n_W_prev,n_C_prev))
dW = np.zeros((f,f,n_C_prev,n_C))
db = np.zeros((f,f,n_C_prev,n_C))
A_prev_pad = zero_pad(A_prev, pad)
dA_prev_pad = zero_pad(dA_prev, pad)

for i in range(m):
    a_prev_pad = A_prev_pad[i,:,:,:]
    da_prev_pad = dA_prev_pad[i,:,:,:]
    for h in range(n_H):           # loop over vertical axis of the output volume
        for w in range(n_W):       # loop over horizontal axis of the output volume
            for c in range(n_C):   # loop over the channels of the output volume

                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f
                a_slice = a_prev_pad[vert_start:vert_end,horiz_start:horiz_end,:]
                da_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :] += W[:, :, :, c] * dZ[i, h, w, c]
                dW[:, :, :, c] += a_slice * dZ[i, h, w, c]
                db[:, :, :, c] += dZ[i, h, w, c]
    dA_prev[i, :, :, :] = da_prev_pad[pad:-pad,pad:-pad,:]
assert(dA_prev.shape == (m, n_H_prev, n_W_prev, n_C_prev))
return dA_prev, dW, db

```

Pooling layer - backward pass

```

def create_mask_from_window(x):
    mask = np.max(x)
    return mask

def distribute_value(dz, shape):
    (n_H, n_W) = shape
    average = dz / (n_H * n_W)
    a = np.ones((n_H,n_W))*average
    return a

def pool_backward(dA, cache, mode = "max"):
    (A_prev, hparameters) = cache
    stride = hparameters['stride']
    f = hparameters['f']

    m, n_H_prev, n_W_prev, n_C_prev = A_prev.shape
    m, n_H, n_W, n_C = dA.shape

```

```

dA_prev = np.zeros(A_prev.shape)

for i in range(m):           # loop over the training examples
    a_prev = A_prev[i,:,:,:]
    for h in range(n_H):      # loop on the vertical axis
        for w in range(n_W):  # loop on the horizontal axis
            for c in range(n_C): # loop over the channels (depth)
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                if mode == "max":
                    a_prev_slice = a_prev[vert_start:vert_end,horiz_start:horiz_end,c]
                    mask = create_mask_from_window(a_prev_slice)
                    dA_prev[i, vert_start: vert_end, horiz_start: horiz_end, c] += np.multiply(mask, dA[i, h,
w, c])

                elif mode == "average":
                    da = dA[i,h,w,c]
                    shape = (f,f)
                    dA_prev[i, vert_start: vert_end, horiz_start: horiz_end, c] += distribute_value(da, shape)
assert(dA_prev.shape == A_prev.shape)
return dA_prev

```

Convolution model - Application

Create placeholders

```

def create_placeholders(n_H0, n_W0, n_C0, n_y):
    X = tf.placeholder(tf.float32, shape=(None, n_H0, n_W0, n_C0))
    Y = tf.placeholder(tf.float32, shape=(None, n_y))
    return X, Y

```

Initialize parameters

```

def initialize_parameters():
    tf.set_random_seed(1)
    W1 = tf.get_variable("W1", [4, 4, 3, 8], initializer=tf.contrib.layers.xavier_initializer(seed=0))
    W2 = tf.get_variable("W2", [2, 2, 8, 16], initializer=tf.contrib.layers.xavier_initializer(seed=0))

    parameters = {"W1": W1,
                  "W2": W2}
    return parameters

```

Forward propagation

```

def forward_propagation(X, parameters):
    W1 = parameters['W1']
    W2 = parameters['W2']

```

```

Z1 = tf.nn.conv2d(X, W1, strides=(1, 1, 1, 1), padding='SAME')
A1 = tf.nn.relu(Z1)
P1 = tf.nn.max_pool(A1, ksize=(1, 8, 8, 1), strides=(1, 8, 8, 1), padding='SAME')
Z2 = tf.nn.conv2d(P1, W2, strides=(1, 1, 1, 1), padding='SAME')
A2 = tf.nn.relu(Z2)
P2 = tf.nn.max_pool(A2, ksize=(1, 4, 4, 1), strides=(1, 4, 4, 1), padding='SAME')
P2 = tf.contrib.layers.flatten(P2)
Z3 = tf.contrib.layers.fully_connected(P2, num_outputs = 6, activation_fn=None)
return Z3

```

Compute cost

```

def compute_cost(Z3, Y):
    cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits_v2(logits = Z3, labels = Y))
    return cost

```

Model

```

def model(X_train, Y_train, X_test, Y_test, learning_rate = 0.009,
          num_epochs = 100, minibatch_size = 64, print_cost = True):
    ops.reset_default_graph()                      # to be able to rerun the model without overwriting tf
variables
    tf.set_random_seed(1)                         # to keep results consistent (tensorflow seed)
    seed = 3                                     # to keep results consistent (numpy seed)
    (m, n_H0, n_W0, n_C0) = X_train.shape
    n_y = Y_train.shape[1]
    costs = []
    X, Y = create_placeholders(n_H0,n_W0,n_C0,n_y)
    parameters = initialize_parameters()
    Z3 = forward_propagation(X,parameters)
    cost = compute_cost(Z3,Y)

```

```
optimizer = tf.train.AdamOptimizer(learning_rate = learning_rate).minimize(cost)
```

```
init = tf.global_variables_initializer()
```

```
with tf.Session() as sess:
```

```
    sess.run(init)
```

```
    for epoch in range(num_epochs):
```

```
        minibatch_cost = 0.
```

```
        num_minibatches = int(m / minibatch_size)
```

```
        seed = seed + 1
```

```
        minibatches = random_mini_batches(X_train, Y_train, minibatch_size, seed)
```

```
        for minibatch in minibatches:
```

```
            (minibatch_X, minibatch_Y) = minibatch
```

```
            _, temp_cost = sess.run([optimizer, cost], feed_dict={X: minibatch_X, Y: minibatch_Y})
```

```

minibatch_cost += temp_cost / num_minibatches

if print_cost == True and epoch % 5 == 0:
    print ("Cost after epoch %i: %f" % (epoch, minibatch_cost))
if print_cost == True and epoch % 1 == 0:
    costs.append(minibatch_cost)

plt.plot(np.squeeze(costs))
plt.ylabel('cost')
plt.xlabel('iterations (per tens)')
plt.title("Learning rate =" + str(learning_rate))
plt.show()
predict_op = tf.argmax(Z3, 1)
correct_prediction = tf.equal(predict_op, tf.argmax(Y, 1))

accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
print(accuracy)
train_accuracy = accuracy.eval({X: X_train, Y: Y_train})
test_accuracy = accuracy.eval({X: X_test, Y: Y_test})
print("Train Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)

return train_accuracy, test_accuracy, parameters

```

结论分析与体会：

卷积层的反向传播相比前向传播复杂一些。先实现 dA, dW, db 的计算，反向传播函数 `conv_backward` 需要把所有的训练样本的过滤器、权值、高度、宽度都要加进来，然后使用千米三个公式计算对应的梯度

前向传播首先是经过卷积层，然后滑动地取卷积层最大值构成了池化层，需要记录最大值的位置，才能反向传播到卷积层。

在 TensorFlow 里面有一些可以直接拿来用的函数，比如

`tf.nn.conv2d`, `tf.nn.relu()`, `tf.contrib.layers.flatten()`, `tf.contrib.layers.fully_connected()`.

前向传播模型的大概步骤就是

CONV2D→RELU→MAXPOOL→CONV2D→RELU→MAXPOOL→FULLCONNECTED